


```
from google.colab import files
uploaded = files.upload()

import pandas as pd

# Load dataset
df = pd.read_csv(next(iter(uploaded)))
df.head()
```

 Choose Files

Housing.csv

- **Housing.csv**(text/csv) - 29981 bytes, last modified: 5/8/2025 - 100% done

Saving Housing.csv to Housing (1).csv

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea	furn
0	13300000	7420	4	2	3	yes	no	no	no	yes	2	yes	
1	12250000	8960	4	4	4	yes	no	no	no	yes	3	no	
2	12250000	9960	3	2	2	yes	no	yes	no	no	2	yes	
3	12215000	7500	4	2	2	yes	no	yes	no	yes	3	yes	
4	11410000	7420	4	1	2	yes	yes	yes	no	yes	2	no	

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
# Step 2: Data Preprocessing
# Display basic info
df.info()

# Handle missing values: fill with median for numerics
df.fillna(df.median(numeric_only=True), inplace=True)

# Get actual column names
print("Columns:", df.columns)

# Define target column and features
target = df.columns[-1] # Last column is the target
X = df.drop(columns=[target])
y = df[target]


# Identify categorical columns
categorical_features = X.select_dtypes(include=['object', 'bool']).columns.tolist()
print("Categorical features:", categorical_features)

# Apply One-Hot Encoding to categorical features
X = pd.get_dummies(X, columns=categorical_features, drop_first=True)

# Scale numeric features only
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X) # Now X contains only numeric data

# Done preprocessing
print("Preprocessing complete. Shape:", X_scaled.shape)
```

 <class 'pandas.core.frame.DataFrame'>

RangeIndex: 545 entries, 0 to 544

Data columns (total 13 columns):

#	Column	Non-Null Count	Dtype
0	price	545 non-null	int64
1	area	545 non-null	int64
2	bedrooms	545 non-null	int64
3	bathrooms	545 non-null	int64
4	stories	545 non-null	int64
5	mainroad	545 non-null	object
6	guestroom	545 non-null	object
7	basement	545 non-null	object
8	hotwaterheating	545 non-null	object
9	airconditioning	545 non-null	object
10	parking	545 non-null	int64
11	prefarea	545 non-null	object

```

12  furnishingstatus  545 non-null    object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
Columns: Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
               'guestroom', 'basement', 'hotwaterheating', 'airconditioning',
               'parking', 'prefarea', 'furnishingstatus'],
              dtype='object')
Categorical features: ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'prefarea']
Preprocessing complete. Shape: (545, 12)

```

Step 3: Select Features and Target

```
target = 'price' # Correct lowercase column name
```

```
X = df.drop(columns=[target])
```

```
y = df[target]
```

Convert categorical features in X using one-hot encoding

```
categorical_cols = X.select_dtypes(include=['object', 'bool']).columns.tolist()
```

```
X = pd.get_dummies(X, columns=categorical_cols, drop_first=True)
```

Scale features

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

Train-test split

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.tree import DecisionTreeRegressor
```

```
from sklearn.metrics import mean_squared_error
```

Linear Regression

```
lr_model = LinearRegression()
```

```
lr_model.fit(X_train, y_train)
```

```
lr_preds = lr_model.predict(X_test)
```

```
lr_mse = mean_squared_error(y_test, lr_preds)
```

Decision Tree Regressor

```
dt_model = DecisionTreeRegressor(random_state=42)
```

```
dt_model.fit(X_train, y_train)
```

```
dt_preds = dt_model.predict(X_test)
```

```
dt_mse = mean_squared_error(y_test, dt_preds)
```

```
print(f"Linear Regression MSE: {lr_mse:.2f}")
```

```
print(f"Decision Tree MSE: {dt_mse:.2f}")
```



```

Linear Regression MSE: 1754318687330.67
Decision Tree MSE: 2642802637614.68

```

Step 5: Plot predictions

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(10, 5))
```

```
plt.plot(y_test.values, label='Actual Prices', color='blue')
```

```
plt.plot(lr_preds, label='Linear Regression Predictions', color='green')
```

```
plt.plot(dt_preds, label='Decision Tree Predictions', color='red')
```

```
plt.title('House Price Prediction')
```

```
plt.xlabel('Sample')
```

```
plt.ylabel('Price')
```

```
plt.legend()
```

```
plt.grid(True)
```

```
plt.show()
```

